

## Session 16 – Assignment

1. Use the below given data set

Data Set

2. Perform the below given activities:

a. Predict the no of comments in next H hrs

Note:-

1. Use LASSO, Elastic Net and Ridge and other regression techniques that are covered in the module

2. Report the training accuracy and test accuracy

3. compare with linear models and report the accuracy

4. create a graph displaying the accuracy of all models

### Attribute Information:

(39 - This describes the H hrs, for which we have the target variable/ comments received.

, 54 -Target Variable - Decimal Target The no of comments in next H hrs(H is given in Feature no 39).

39

H Local

Decimal(0-23) Encoding

Other feature

This describes the H hrs, for which we have the target variable/ comments received.

54

Target Variable

Decimal

Target

The no of comments in next H hrs(H is given in Feature no 39).

## Prediction Accuracy

*A good learner is the one which has good prediction accuracy; in other words, which has the smallest prediction error.*

Let us try to understand the prediction problem intuitively. Consider the simple case of fitting a linear regression model to the observed data. A model is a good fit, if it provides a high  $R^2$  value.

1. Use LASSO, Elastic Net and Ridge and other regression techniques that are covered in the module

```
library(tidyverse)
```

```
library(caret)
```

```
library(glmnet)
```

```
# Load the data
```

```
setwd("~/Dataset/Dataset/Training")
```

```
Features_Variant_1 <-
```

```
read.csv("C:/users/seshan/Documents/Dataset/Dataset/Training/Features_Variant_1.csv")
```

```
View(Features_Variant_1)
```

```
Features.data <- na.omit(Features_Variant_1)
```

```
# Split the data into training and test set
```

```
set.seed(123)
```

```
training.samples <- Features$X0.19 %>%
```

```
  createDataPartition(p = 0.8, list = FALSE)
```

```
train.data <- Features_Variant_1[training.samples, ]
```

```
test.data <- Features_Variant_1[-training.samples, ]
```

```
# Predictor variables
```

```
x <- model.matrix(X0.19~., train.data)[-1]
```

```
# Outcome variable
```

```
y <- train.data$X0.19
```

```
glmnet(x, y, alpha = 1, lambda = NULL)
```

```
# Find the best lambda using cross-validation
```

```
set.seed(123)
```

```
cv <- cv.glmnet(x, y, alpha = 0)
```

```
# Display the best lambda value
```

```

cv$lambda.min

plot(cv$lambda.min)

# Fit the final model on the training data

model <- glmnet(x, y, alpha = 0, lambda = cv$lambda.min)

plot(model)

# Display regression coefficients

coef(model)


# Make predictions on the test data

x.test <- model.matrix(X0.19 ~., test.data)[-1]

predictions <- model %>% predict(x.test) %>% as.vector()

# Model performance metrics

data.frame(

  RMSE = RMSE(predictions, test.data$X0.19),

  Rsquare = R2(predictions, test.data$X0.19)

)


#Computing lasso regression

# Find the best lambda using cross-validation

set.seed(123)

cv <- cv.glmnet(x, y, alpha = 1)

# Display the best lambda value

cv$lambda.min

# Fit the final model on the training data

model <- glmnet(x, y, alpha = 1, lambda = cv$lambda.min)

```

```

# Display regression coefficients
coef(model)

# Make predictions on the test data
x.test <- model.matrix(X0.19 ~., test.data)[-1]

predictions <- model %>% predict(x.test) %>% as.vector()

# Model performance metrics
data.frame(

  RMSE = RMSE(predictions, test.data$X0.19),

  Rsquare = R2(predictions, test.data$X0.19))

# Computing elastic net regression

# Build the model using the training set
set.seed(123)

model <- train(X0.19 ~., data = train.data, method = "glmnet",

  trControl = trainControl("cv", number = 10),

  tuneLength = 10

)

# Best tuning parameter
model$bestTune

plot(model$bestTune)

# Coefficient of the final model. You need

# to specify the best lambda
coef(model$finalModel, model$bestTune$lambda)

# Make predictions on the test data
x.test <- model.matrix(X0.19 ~., test.data)[-1]

predictions <- model %>% predict(x.test)

```

```

# Model performance metrics

data.frame(

  RMSE = RMSE(predictions, test.data$X0.19),

  Rsquare = R2(predictions, test.data$X0.19)

)

#Comparing the different models

#Using caret package

#Setup a grid range of lambda values:

lambda <- 10^seq(-3, 3, length = 100)

#Compute ridge regression

# Build the model

set.seed(123)

ridge <- train(

  X0.19 ~., data = train.data, method = "glmnet",

  trControl = trainControl("cv", number = 10),

  tuneGrid = expand.grid(alpha = 0, lambda = lambda)

)

# Model coefficients

coef(ridge$finalModel, ridge$bestTune$lambda)

# Make predictions

predictions <- ridge %>% predict(test.data)

plot(predictions)

# Model prediction performance

data.frame(

  RMSE = RMSE(predictions, test.data$X0.19),

```

```

Rsquare = R2(predictions, test.data$X0.19)
)

#Compute lasso regression

# Build the model

set.seed(123)

lasso <- train(
  X0.19 ~., data = train.data, method = "glmnet",
  trControl = trainControl("cv", number = 10),
  tuneGrid = expand.grid(alpha = 1, lambda = lambda)
)

plot( lasso)

# Model coefficients

coef(lasso$finalModel, lasso$bestTune$lambda)

# Make predictions

predictions <- lasso %>% predict(test.data)

# Model prediction performance

data.frame(
  RMSE = RMSE(predictions, test.data$X0.19),
  Rsquare = R2(predictions, test.data$X0.19)
)

#Elastic net regression

# Build the model

set.seed(123)

elastic <- train(

```

```

X0.19 ~., data = train.data, method = "glmnet",
trControl = trainControl("cv", number = 10),
tuneLength = 10
)

# Model coefficients
coef(elastic$finalModel, elastic$bestTune$lambda)


# Make predictions
predictions <- elastic %>% predict(test.data)
plot( predictions)

# Model prediction performance
data.frame(
  RMSE = RMSE(predictions, test.data$X0.19),
  Rsquare = R2(predictions, test.data$X0.19)
)

#Comparing models performance:
models <- list(ridge = ridge, lasso = lasso, elastic = elastic)
resamples(models) %>% summary( metric = "RMSE")


#k-fold Cross Validation

# load the library
library(caret)

# define training control

```

```

train_control <- trainControl(method="cv", number=10)

# fix the parameters of the algorithm

grid <- expand.grid(.fl=c(0), .usekernel=c(FALSE))

# train the model

model <- train(X0.19~., data=Features_Variant_1, trControl=train_control, method="nb",
tuneGrid=grid)

# summarize results

print(model)


# load the library

library(caret)


# define training control

train_control <- trainControl(method="repeatedcv", number=10, repeats=3)

# train the model

model <- train(X0.19~., data=Features_Variant_1, trControl=train_control, method="nb")

# summarize results

print(model)


#create a graph displaying the accuracy of all models

plot(model)

plot(varImp(ridge$finalModel))

plot(cv)

plot(ridge)

hist(Features$X0.19,col = "green")

hist(Features$X24,col = "red")

```



```
hist(Features$X11.291044776119403,col = 'yellow')
```

```
fit = glmnet(x, y)
```

```
plot(fit)
```

```
cvfit = cv.glmnet(x, y)
```

```
plot(cvfit)
```

```
tfit=glmnet(x,y,lower=-.7,upper=.5)
```

```
plot(tfit)
```

```
> library(tidyverse)
> library(caret)
> library(glmnet)
> # Load the data
> setwd("~/Dataset/Dataset/Training")
>
> Features_Variant_1 <- read.csv("C:/users/seshan/Documents/Dataset/Dataset/Training/
Features_Variant_1.csv")
> View(Features_Variant_1)
> Features.data <- na.omit(Features_Variant_1)
> # Split the data into training and test set
> set.seed(123)
> training.samples <- Features$X0.19 %>%
+   createDataPartition(p = 0.8, list = FALSE)
> train.data <- Features_Variant_1[training.samples, ]
> test.data <- Features_Variant_1[-training.samples, ]
> # Predictor variables
> x <- model.matrix(X0.19~., train.data)[,-1]
> # Outcome variable
> y <- train.data$X0.19
> glmnet(x, y, alpha = 1, lambda = NULL)
```

```
Call:  glmnet(x = x, y = y, alpha = 1, lambda = NULL)
```

	Df	%Dev	Lambda
[1,]	0	0.00000	18.250000
[2,]	1	0.04881	16.630000
[3,]	1	0.08933	15.150000
[4,]	1	0.12300	13.810000
[5,]	1	0.15090	12.580000
[6,]	1	0.17410	11.460000
[7,]	1	0.19330	10.440000
[8,]	1	0.20930	9.516000
[9,]	1	0.22260	8.671000
[10,]	1	0.23360	7.900000
[11,]	1	0.24270	7.198000

[12,]	1	0.25030	6.559000
[13,]	1	0.25660	5.976000
[14,]	2	0.26350	5.445000
[15,]	2	0.26950	4.962000
[16,]	3	0.27660	4.521000
[17,]	3	0.28290	4.119000
[18,]	4	0.28840	3.753000
[19,]	4	0.29290	3.420000
[20,]	4	0.29670	3.116000
[21,]	6	0.30040	2.839000
[22,]	6	0.30370	2.587000
[23,]	6	0.30640	2.357000
[24,]	7	0.30870	2.148000
[25,]	7	0.31060	1.957000
[26,]	8	0.31210	1.783000
[27,]	8	0.31350	1.625000
[28,]	8	0.31490	1.480000
[29,]	8	0.31610	1.349000
[30,]	8	0.31710	1.229000
[31,]	8	0.31790	1.120000
[32,]	8	0.31860	1.020000
[33,]	8	0.31920	0.929700
[34,]	8	0.31970	0.847100
[35,]	9	0.32020	0.771900
[36,]	9	0.32060	0.703300
[37,]	9	0.32090	0.640800
[38,]	9	0.32120	0.583900
[39,]	10	0.32150	0.532000
[40,]	10	0.32180	0.484800
[41,]	11	0.32200	0.441700
[42,]	11	0.32240	0.402500
[43,]	11	0.32270	0.366700
[44,]	12	0.32290	0.334100
[45,]	12	0.32320	0.304400
[46,]	13	0.32330	0.277400
[47,]	14	0.32370	0.252800
[48,]	16	0.32390	0.230300
[49,]	17	0.32420	0.209800
[50,]	17	0.32440	0.191200
[51,]	18	0.32460	0.174200
[52,]	22	0.32480	0.158700
[53,]	22	0.32500	0.144600
[54,]	23	0.32520	0.131800
[55,]	24	0.32540	0.120100
[56,]	24	0.32550	0.109400
[57,]	26	0.32570	0.099690
[58,]	29	0.32590	0.090830
[59,]	29	0.32600	0.082770
[60,]	30	0.32620	0.075410
[61,]	30	0.32630	0.068710
[62,]	30	0.32640	0.062610
[63,]	31	0.32640	0.057050
[64,]	31	0.32650	0.051980
[65,]	34	0.32660	0.047360
[66,]	35	0.32660	0.043150
[67,]	36	0.32680	0.039320
[68,]	36	0.32700	0.035830
[69,]	37	0.32730	0.032640
[70,]	38	0.32750	0.029740
[71,]	37	0.32770	0.027100
[72,]	40	0.32790	0.024690
[73,]	40	0.32820	0.022500
[74,]	45	0.32840	0.020500
[75,]	45	0.32880	0.018680

```

[76,] 46 0.32930 0.017020
[77,] 46 0.32960 0.015510
[78,] 47 0.32980 0.014130
[79,] 47 0.33000 0.012880
[80,] 47 0.33010 0.011730
[81,] 47 0.33030 0.010690
[82,] 47 0.33040 0.009740
[83,] 47 0.33050 0.008875
[84,] 47 0.33060 0.008086
[85,] 48 0.33070 0.007368
[86,] 48 0.33080 0.006713
[87,] 48 0.33090 0.006117
[88,] 48 0.33090 0.005574
[89,] 48 0.33100 0.005078
[90,] 48 0.33100 0.004627
[91,] 48 0.33110 0.004216
[92,] 48 0.33110 0.003842
[93,] 49 0.33120 0.003500
[94,] 49 0.33120 0.003189
[95,] 49 0.33120 0.002906
[96,] 49 0.33120 0.002648
[97,] 49 0.33130 0.002413
[98,] 49 0.33130 0.002198
[99,] 49 0.33130 0.002003
[100,] 49 0.33130 0.001825
> # Find the best lambda using cross-validation
> set.seed(123)
> cv <- cv.glmnet(x, y, alpha = 0)
> # Display the best lambda value
> cv$lambda.min
[1] 3.189382
> plot(cv$lambda.min)
> # Fit the final model on the training data
> model <- glmnet(x, y, alpha = 0, lambda = cv$lambda.min)
> plot(model)
> # Display regression coefficients
> coef(model)
54 x 1 sparse Matrix of class "dgCMatrix"
      s0
(Intercept) -2.538449e+00
x634995      4.086976e-08
x0           -2.391902e-05
x463         -1.118531e-05
x1           7.735292e-04
x0.0         1.350121e-02
x806.0       1.553809e-03
x11.291044776119403 9.304117e-03
x1.0         1.793778e-04
x70.49513846124168 2.001045e-02
x0.0.1       -1.965736e-02
x806.0.1     -1.021345e-03
x7.574626865671642 2.932718e-02
x0.0.2       1.122924e-01
x69.435826365571 -3.244738e-03
x0.0.3       -7.678915e-03
x76.0        -2.361222e-04
x2.6044776119402986 2.333540e-02
x0.0.4       1.156420e-01
x8.50550186882253 4.524854e-03
x0.0.5       -1.178544e-02
x806.0.2     -1.639273e-03
x10.649253731343284 -3.665694e-03
x1.0.1       -1.572424e-02
x70.25478763764251 5.195732e-03

```

```

x.69.0      1.960444e-04
x806.0.3    -5.160122e-04
x4.970149253731344  4.298933e-02
x0.0.6      -3.755897e-02
x69.85058043098057  4.831630e-03
x0.1        3.362289e-03
x0.2        1.151692e-01
x0.3        2.359120e-02
x0.4        5.274079e-04
x0.5        6.427015e-02
x65         -1.912346e-01
x166        2.483775e-04
x2          1.745596e-03
x0.6        .
x24         3.970945e-01
x0.7        -8.263988e-01
x0.8        -4.301798e-01
x0.9        -4.379730e-01
x1.1        7.482821e-01
x0.10       4.288169e-01
x0.11       5.881030e-01
x0.12       -2.262055e-01
x0.13       -5.948819e-01
x0.14       4.781008e-01
x0.15       -7.372765e-02
x0.16       1.043352e+00
x0.17       -1.065425e-01
x0.18       -2.291072e-01
x1.2       -5.230538e-01
>
> # Make predictions on the test data
> x.test <- model.matrix(x0.19 ~., test.data)[,-1]
> predictions <- model %>% predict(x.test) %>% as.vector()
> # Model performance metrics
> data.frame(
+   RMSE = RMSE(predictions, test.data$x0.19),
+   Rsquare = R2(predictions, test.data$x0.19)
+ )
      RMSE   Rsquare
1 34.21296 0.3043655
>
> #Computing lasso regression
> # Find the best lambda using cross-validation
> set.seed(123)
> cv <- cv.glmnet(x, y, alpha = 1)
> # Display the best lambda value
> cv$lambda.min
[1] 0.7033004
> # Fit the final model on the training data
> model <- glmnet(x, y, alpha = 1, lambda = cv$lambda.min)
> # Display regression coefficients
> coef(model)
54 x 1 sparse Matrix of class "dgCMatrix"
              s0
(Intercept)  5.0667129773
x634995      .
x0           .
x463         .
x1           .
x0.0         .
x806.0       .
x11.291044776119403 .
x1.0         .
x70.49513846124168 0.0131953053

```

```

X0.0.1 .
X806.0.1 .
X7.574626865671642 0.0278254674
X0.0.2 0.0955320465
X69.435826365571 .
X0.0.3 .
X76.0 .
X2.6044776119402986 .
X0.0.4 0.1082157028
X8.50550186882253 .
X0.0.5 .
X806.0.2 .
X10.649253731343284 .
X1.0.1 .
X70.25478763764251 .
X.69.0 .
X806.0.3 .
X4.970149253731344 .
X0.0.6 .
X69.85058043098057 .
X0.1 .
X0.2 0.1661874069
X0.3 .
X0.4 .
X0.5 0.0264102499
X65 -0.1704729896
X166 .
X2 0.0007766533
X0.6 .
X24 0.0658390632
X0.7 .
X0.8 .
X0.9 .
X1.1 .
X0.10 .
X0.11 .
X0.12 .
X0.13 .
X0.14 .
X0.15 .
X0.16 .
X0.17 .
X0.18 .
X1.2 .
> # Make predictions on the test data
> x.test <- model.matrix(X0.19 ~., test.data)[,-1]
> predictions <- model %>% predict(x.test) %>% as.vector()
> # Model performance metrics
> data.frame(
+   RMSE = RMSE(predictions, test.data$X0.19),
+   Rsquare = R2(predictions, test.data$X0.19))
  RMSE Rsquare
1 34.35851 0.2981491
> #Computing elastic net regression
> # Build the model using the training set
> set.seed(123)
> model <- train(X0.19 ~., data = train.data, method = "glmnet",
+   trControl = trainControl("cv", number = 10),
+   tuneLength = 10
+ )
> # Best tuning parameter
> model$bestTun
  alpha lambda
96    1 0.5547923

```

```

> plot(model$bestTun)
> # Coefficient of the final model. You need
> # to specify the best lambda
> coef(model$finalModel, model$bestTune$lambda)
54 x 1 sparse Matrix of class "dgCMatrix"
      1
(Intercept)      3.243771e+00
x634995          .
x0             -1.451030e-06
x463            .
x1              .
x0.0            .
x806.0          .
x11.291044776119403 .
x1.0            .
x70.49513846124168 1.537445e-02
x0.0.1          .
x806.0.1        .
x7.574626865671642 2.330706e-02
x0.0.2          1.020902e-01
x69.435826365571  .
x0.0.3          .
x76.0           .
x2.6044776119402986 .
x0.0.4          1.150414e-01
x8.50550186882253  .
x0.0.5          .
x806.0.2        .
x10.649253731343284 .
x1.0.1          .
x70.25478763764251 .
x.69.0          .
x806.0.3        .
x4.970149253731344 .
x0.0.6          .
x69.85058043098057 .
x0.1            .
x0.2            1.650641e-01
x0.3            .
x0.4            .
x0.5            2.819936e-02
x65             -1.762485e-01
x166            .
x2              9.079099e-04
x0.6            .
x24             1.461588e-01
x0.7            .
x0.8            .
x0.9            .
x1.1            .
x0.10           .
x0.11           .
x0.12           .
x0.13           .
x0.14           .
x0.15           .
x0.16           .
x0.17           .
x0.18           .
x1.2            .
> # Make predictions on the test data
> x.test <- model.matrix(x0.19 ~., test.data)[,-1]
> predictions <- model %>% predict(x.test)
> # Model performance metrics

```

```

> data.frame(
+   RMSE = RMSE(predictions, test.data$X0.19),
+   Rsquare = R2(predictions, test.data$X0.19)
+ )
      RMSE   Rsquare
1 34.31238 0.2993523
> #Comparing the different models
> #Using caret package
> #Setup a grid range of lambda values:
>   lambda <- 10^seq(-3, 3, length = 100)
> #Compute ridge regression
>   # Build the model
>   set.seed(123)
>   ridge <- train(
+     X0.19 ~., data = train.data, method = "glmnet",
+     trControl = trainControl("cv", number = 10),
+     tuneGrid = expand.grid(alpha = 0, lambda = lambda)
+   )
>   # Model coefficients
>   coef(ridge$finalModel, ridge$bestTune$lambda)
54 x 1 sparse Matrix of class "dgCMatrix"
      1
(Intercept)      -2.780069e+00
X634995          5.350956e-08
X0              -2.364834e-05
X463            -1.329873e-05
X1              6.216835e-04
X0.0            2.038074e-02
X806.0          1.700351e-03
X11.291044776119403 8.550170e-03
X1.0            -4.004515e-03
X70.49513846124168 2.306081e-02
X0.0.1          -5.381912e-02
X806.0.1        -1.655099e-03
X7.574626865671642 2.994944e-02
X0.0.2          1.245636e-01
X69.435826365571 -7.884156e-03
X0.0.3          2.622963e-02
X76.0           -3.747457e-04
X2.6044776119402986 2.978466e-02
X0.0.4          1.358288e-01
X8.50550186882253 3.211969e-03
X0.0.5          -1.596299e-02
X806.0.2        -2.000906e-03
X10.649253731343284 -3.364991e-03
X1.0.1          -1.984096e-02
X70.25478763764251 7.880343e-03
X.69.0          -2.562913e-04
X806.0.3        9.298950e-05
X4.970149253731344 5.279758e-02
X0.0.6          -3.668527e-02
X69.85058043098057 3.578648e-03
X0.1            2.690513e-03
X0.2            1.241846e-01
X0.3            3.087028e-02
X0.4            -5.186126e-03
X0.5            6.542507e-02
X65             -1.911067e-01
X166            2.396281e-04
X2              1.857963e-03
X0.6            .
X24             4.067486e-01
X0.7            -8.764705e-01
X0.8            -4.849690e-01

```

```

X0.9 -4.740900e-01
X1.1 7.672004e-01
X0.10 4.799051e-01
X0.11 6.261735e-01
X0.12 -2.236307e-01
X0.13 -6.182879e-01
X0.14 5.103142e-01
X0.15 -3.737265e-02
X0.16 1.097674e+00
X0.17 -1.186700e-01
X0.18 -2.627125e-01
X1.2 -5.621735e-01

```

```

> # Make predictions
> predictions <- ridge %>% predict(test.data)
> plot(predictions)
> # Model prediction performance
> data.frame(
+   RMSE = RMSE(predictions, test.data$X0.19),
+   Rsquare = R2(predictions, test.data$X0.19)
+ )

```

```

      RMSE   Rsquare
1 34.15522 0.3061988

```

```

>
> #Compute lasso regression
> # Build the model
> set.seed(123)
> lasso <- train(
+   X0.19 ~., data = train.data, method = "glmnet",
+   trControl = trainControl("cv", number = 10),
+   tuneGrid = expand.grid(alpha = 1, lambda = lambda)
+ )

```

In nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :  
There were missing values in resampled performance measures.

```

> plot(lasso)
> # Model coefficients
> coef(lasso$finalModel, lasso$bestTune$lambda)

```

54 x 1 sparse Matrix of class "dgCMatrix"

```

      1
(Intercept) 2.988187e+00
x634995      .
X0          -2.504262e-06
x463        .
x1           .
X0.0         .
x806.0       .
x11.291044776119403 .
x1.0         .
x70.49513846124168 1.574129e-02
X0.0.1       .
x806.0.1     .
x7.574626865671642 2.257750e-02
X0.0.2       1.029959e-01
x69.435826365571 .
X0.0.3       .
x76.0        .
x2.6044776119402986 .
X0.0.4       1.160087e-01
x8.50550186882253 .
X0.0.5       .
x806.0.2     .
x10.649253731343284 .
x1.0.1       .
x70.25478763764251 .

```



```

X.69.0      .
X806.0.3    .
X4.970149253731344 .
X0.0.6      .
X69.85058043098057 .
X0.1        .
X0.2        1.649295e-01
X0.3        .
X0.4        .
X0.5        2.844279e-02
X65         -1.770593e-01
X166        .
X2          9.269250e-04
X0.6        .
X24         1.574830e-01
X0.7        .
X0.8        .
X0.9        .
X1.1        .
X0.10       .
X0.11       .
X0.12       .
X0.13       .
X0.14       .
X0.15       .
X0.16       .
X0.17       .
X0.18       .
X1.2        .
> # Make predictions
> predictions <- lasso %>% predict(test.data)
> # Model prediction performance
> data.frame(
+   RMSE = RMSE(predictions, test.data$X0.19),
+   Rsquare = R2(predictions, test.data$X0.19)
+ )
  RMSE   Rsquare
1 34.3056 0.2995397
> #Elastic net regression
> # Build the model
> set.seed(123)
> elastic <- train(
+   X0.19 ~., data = train.data, method = "glmnet",
+   trControl = trainControl("cv", number = 10),
+   tuneLength = 10
+ )
> # Model coefficients
> coef(elastic$finalModel, elastic$bestTune$lambda)
54 x 1 sparse Matrix of class "dgCMatrix"
      1
(Intercept) 3.243771e+00
X634995     .
X0          -1.451030e-06
X463        .
X1          .
X0.0        .
X806.0      .
X11.291044776119403 .
X1.0        .
X70.49513846124168 1.537445e-02
X0.0.1      .
X806.0.1    .
X7.574626865671642 2.330706e-02
X0.0.2      1.020902e-01

```

```

X69.435826365571 .
X0.0.3 .
X76.0 .
X2.6044776119402986 .
X0.0.4 1.150414e-01
X8.50550186882253 .
X0.0.5 .
X806.0.2 .
X10.649253731343284 .
X1.0.1 .
X70.25478763764251 .
X.69.0 .
X806.0.3 .
X4.970149253731344 .
X0.0.6 .
X69.85058043098057 .
X0.1 .
X0.2 1.650641e-01
X0.3 .
X0.4 .
X0.5 2.819936e-02
X65 -1.762485e-01
X166 .
X2 9.079099e-04
X0.6 .
X24 1.461588e-01
X0.7 .
X0.8 .
X0.9 .
X1.1 .
X0.10 .
X0.11 .
X0.12 .
X0.13 .
X0.14 .
X0.15 .
X0.16 .
X0.17 .
X0.18 .
X1.2 .
>
> # Make predictions
> predictions <- elastic %>% predict(test.data)
> plot( predictions)
> # Model prediction performance
> data.frame(
+   RMSE = RMSE(predictions, test.data$X0.19),
+   Rsquare = R2(predictions, test.data$X0.19)
+ )
  RMSE Rsquare
1 34.31238 0.2993523
> #Comparing models performance:
> models <- list(ridge = ridge, lasso = lasso, elastic = elastic)
> resamples(models) %>% summary( metric = "RMSE")

Call:
summary.resamples(object = ., metric = "RMSE")

Models: ridge, lasso, elastic
Number of resamples: 10

RMSE
ridge   Min.   1st Qu.   Median   Mean   3rd Qu.   Max. NA's
 21.64890 25.52171 28.01852 27.99756 30.88915 35.66276    0

```

```
lasso 21.52017 25.58002 28.07640 27.93523 30.71819 35.53640 0
elastic 21.51814 25.58351 28.07020 27.93467 30.72364 35.53586 0
```

```
>
>
> #k-fold Cross Validation
> # load the library
> library(caret)
>
> # define training control
> train_control <- trainControl(method="cv", number=10)
> # fix the parameters of the algorithm
> grid <- expand.grid(.fL=c(0), .usekernel=c(FALSE))
> # train the model
> model <- train(X0.19~., data=Features_Variant_1, trControl=train_control, method=
"nb", tuneGrid=grid)
> # summarize results
> print(model)
```

glmnet

32760 samples  
53 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 29484, 29485, 29484, 29484, 29483, 29484, ...

Resampling results across tuning parameters:

alpha	lambda	RMSE	Rsquared	MAE
0.1	0.008432348	28.00932	0.3176368	8.073572
0.1	0.019479818	28.00913	0.3176396	8.072474
0.1	0.045000908	28.00646	0.3176551	8.050171
0.1	0.103957936	28.00846	0.3174455	8.016234
0.1	0.240156321	28.00059	0.3176130	7.975161
0.1	0.554792263	27.98067	0.3181937	7.923823
0.1	1.281642117	27.96815	0.3182085	7.868204
0.1	2.960759592	27.97201	0.3179401	7.737538
0.1	6.839738840	28.04120	0.3164505	7.526685
0.1	15.800684229	28.33270	0.3138865	7.217222
0.2	0.008432348	28.01765	0.3172223	8.072621
0.2	0.019479818	28.01210	0.3174223	8.061405
0.2	0.045000908	28.00976	0.3174643	8.039239
0.2	0.103957936	28.00648	0.3174175	7.996145
0.2	0.240156321	27.98642	0.3180788	7.943594
0.2	0.554792263	27.96948	0.3183579	7.887480
0.2	1.281642117	27.95241	0.3186087	7.775205
0.2	2.960759592	27.96450	0.3182276	7.562917
0.2	6.839738840	28.09348	0.3170458	7.222323
0.2	15.800684229	28.65765	0.3151025	7.133642
0.3	0.008432348	28.00299	0.3177588	8.065477
0.3	0.019479818	28.00437	0.3177326	8.057176
0.3	0.045000908	28.01072	0.3174032	8.031060
0.3	0.103957936	28.00038	0.3176307	7.975466
0.3	0.240156321	27.97652	0.3183974	7.915900
0.3	0.554792263	27.95698	0.3186058	7.846411
0.3	1.281642117	27.94976	0.3185072	7.683206
0.3	2.960759592	27.97534	0.3181303	7.400328
0.3	6.839738840	28.17680	0.3172627	6.985857
0.3	15.800684229	29.13696	0.3103036	7.570747
0.4	0.008432348	28.01830	0.3172446	8.079680
0.4	0.019479818	28.00882	0.3176078	8.058608
0.4	0.045000908	28.01072	0.3173242	8.019561
0.4	0.103957936	27.99402	0.3178502	7.957537
0.4	0.240156321	27.96938	0.3185741	7.892792

0.4	0.554792263	27.94835	0.3187576	7.805517
0.4	1.281642117	27.94390	0.3185641	7.599557
0.4	2.960759592	27.99377	0.3179287	7.243370
0.4	6.839738840	28.28988	0.3165676	6.848175
0.4	15.800684229	29.67198	0.2994897	7.987755
0.5	0.008432348	28.01315	0.3174913	8.081913
0.5	0.019479818	28.00741	0.3176329	8.055975
0.5	0.045000908	28.00761	0.3174288	8.007801
0.5	0.103957936	27.98665	0.3181211	7.942926
0.5	0.240156321	27.96218	0.3187504	7.873342
0.5	0.554792263	27.94327	0.3188512	7.760826
0.5	1.281642117	27.94205	0.3185650	7.519772
0.5	2.960759592	28.02238	0.3174298	7.096167
0.5	6.839738840	28.45004	0.3134711	6.885703
0.5	15.800684229	30.10478	0.2911250	8.323360
0.6	0.008432348	28.01123	0.3175588	8.077084
0.6	0.019479818	28.00535	0.3177547	8.051417
0.6	0.045000908	28.00633	0.3174460	7.995543
0.6	0.103957936	27.98081	0.3183217	7.929786
0.6	0.240156321	27.95596	0.3188777	7.854967
0.6	0.554792263	27.94107	0.3188351	7.718358
0.6	1.281642117	27.94808	0.3183242	7.443956
0.6	2.960759592	28.05422	0.3169109	6.963000
0.6	6.839738840	28.64667	0.3078443	7.060181
0.6	15.800684229	30.45088	0.2909322	8.582636
0.7	0.008432348	28.01327	0.3174691	8.070577
0.7	0.019479818	28.00938	0.3175657	8.046031
0.7	0.045000908	28.00293	0.3175654	7.982522
0.7	0.103957936	27.97610	0.3184735	7.917651
0.7	0.240156321	27.95104	0.3189454	7.835833
0.7	0.554792263	27.94103	0.3187177	7.678114
0.7	1.281642117	27.95454	0.3181156	7.369797
0.7	2.960759592	28.09005	0.3163118	6.846891
0.7	6.839738840	28.85612	0.3001402	7.277639
0.7	15.800684229	30.85446	0.2909322	8.878579
0.8	0.008432348	28.01238	0.3175718	8.070449
0.8	0.019479818	28.00936	0.3175163	8.039224
0.8	0.045000908	28.00020	0.3176370	7.972411
0.8	0.103957936	27.97254	0.3185722	7.906252
0.8	0.240156321	27.94734	0.3189795	7.817664
0.8	0.554792263	27.94076	0.3186304	7.640068
0.8	1.281642117	27.96246	0.3178774	7.295586
0.8	2.960759592	28.13781	0.3151418	6.755044
0.8	6.839738840	29.02499	0.2946305	7.449163
0.8	15.800684229	31.33025	0.2909322	9.216728
0.9	0.008432348	28.01155	0.3175817	8.069278
0.9	0.019479818	28.01019	0.3174714	8.033794
0.9	0.045000908	27.99737	0.3177359	7.963507
0.9	0.103957936	27.96951	0.3186458	7.895521
0.9	0.240156321	27.94458	0.3189922	7.800738
0.9	0.554792263	27.93743	0.3186621	7.602655
0.9	1.281642117	27.97328	0.3175364	7.223312
0.9	2.960759592	28.19981	0.3131783	6.691255
0.9	6.839738840	29.16270	0.2914991	7.572163
0.9	15.800684229	31.88869	0.2909322	9.602631
1.0	0.008432348	28.00984	0.3176179	8.064194
1.0	0.019479818	28.01048	0.3174197	8.027480
1.0	0.045000908	27.99330	0.3178932	7.955851
1.0	0.103957936	27.96647	0.3187201	7.885657
1.0	0.240156321	27.94291	0.3189672	7.784590
1.0	0.554792263	27.93467	0.3186856	7.565986
1.0	1.281642117	27.98635	0.3171131	7.153438
1.0	2.960759592	28.26663	0.3108182	6.661504
1.0	6.839738840	29.27235	0.2909322	7.666803

```
1.0      15.800684229  32.54609  0.2909322  10.044511
```

RMSE was used to select the optimal model using the smallest value.  
The final values used for the model were alpha = 1 and lambda = 0.5547923.

```
>
> # load the library
> library(caret)
>
> # define training control
> train_control <- trainControl(method="repeatedcv", number=10, repeats=3)
> # train the model
> model <- train(X0.19~., data=Features_Variant_1, trControl=train_control, method=
"nb")
> # summarize results
> print(model)
glmnet
```

```
32760 samples
  53 predictor
```

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 29484, 29485, 29484, 29484, 29483, 29484, ...

Resampling results across tuning parameters:

alpha	lambda	RMSE	Rsquared	MAE
0.1	0.008432348	28.00932	0.3176368	8.073572
0.1	0.019479818	28.00913	0.3176396	8.072474
0.1	0.045000908	28.00646	0.3176551	8.050171
0.1	0.103957936	28.00846	0.3174455	8.016234
0.1	0.240156321	28.00059	0.3176130	7.975161
0.1	0.554792263	27.98067	0.3181937	7.923823
0.1	1.281642117	27.96815	0.3182085	7.868204
0.1	2.960759592	27.97201	0.3179401	7.737538
0.1	6.839738840	28.04120	0.3164505	7.526685
0.1	15.800684229	28.33270	0.3138865	7.217222
0.2	0.008432348	28.01765	0.3172223	8.072621
0.2	0.019479818	28.01210	0.3174223	8.061405
0.2	0.045000908	28.00976	0.3174643	8.039239
0.2	0.103957936	28.00648	0.3174175	7.996145
0.2	0.240156321	27.98642	0.3180788	7.943594
0.2	0.554792263	27.96948	0.3183579	7.887480
0.2	1.281642117	27.95241	0.3186087	7.775205
0.2	2.960759592	27.96450	0.3182276	7.562917
0.2	6.839738840	28.09348	0.3170458	7.222323
0.2	15.800684229	28.65765	0.3151025	7.133642
0.3	0.008432348	28.00299	0.3177588	8.065477
0.3	0.019479818	28.00437	0.3177326	8.057176
0.3	0.045000908	28.01072	0.3174032	8.031060
0.3	0.103957936	28.00038	0.3176307	7.975466
0.3	0.240156321	27.97652	0.3183974	7.915900
0.3	0.554792263	27.95698	0.3186058	7.846411
0.3	1.281642117	27.94976	0.3185072	7.683206
0.3	2.960759592	27.97534	0.3181303	7.400328
0.3	6.839738840	28.17680	0.3172627	6.985857
0.3	15.800684229	29.13696	0.3103036	7.570747
0.4	0.008432348	28.01830	0.3172446	8.079680
0.4	0.019479818	28.00882	0.3176078	8.058608
0.4	0.045000908	28.01072	0.3173242	8.019561
0.4	0.103957936	27.99402	0.3178502	7.957537
0.4	0.240156321	27.96938	0.3185741	7.892792
0.4	0.554792263	27.94835	0.3187576	7.805517
0.4	1.281642117	27.94390	0.3185641	7.599557
0.4	2.960759592	27.99377	0.3179287	7.243370

0.4	6.839738840	28.28988	0.3165676	6.848175
0.4	15.800684229	29.67198	0.2994897	7.987755
0.5	0.008432348	28.01315	0.3174913	8.081913
0.5	0.019479818	28.00741	0.3176329	8.055975
0.5	0.045000908	28.00761	0.3174288	8.007801
0.5	0.103957936	27.98665	0.3181211	7.942926
0.5	0.240156321	27.96218	0.3187504	7.873342
0.5	0.554792263	27.94327	0.3188512	7.760826
0.5	1.281642117	27.94205	0.3185650	7.519772
0.5	2.960759592	28.02238	0.3174298	7.096167
0.5	6.839738840	28.45004	0.3134711	6.885703
0.5	15.800684229	30.10478	0.2911250	8.323360
0.6	0.008432348	28.01123	0.3175588	8.077084
0.6	0.019479818	28.00535	0.3177547	8.051417
0.6	0.045000908	28.00633	0.3174460	7.995543
0.6	0.103957936	27.98081	0.3183217	7.929786
0.6	0.240156321	27.95596	0.3188777	7.854967
0.6	0.554792263	27.94107	0.3188351	7.718358
0.6	1.281642117	27.94808	0.3183242	7.443956
0.6	2.960759592	28.05422	0.3169109	6.963000
0.6	6.839738840	28.64667	0.3078443	7.060181
0.6	15.800684229	30.45088	0.2909322	8.582636
0.7	0.008432348	28.01327	0.3174691	8.070577
0.7	0.019479818	28.00938	0.3175657	8.046031
0.7	0.045000908	28.00293	0.3175654	7.982522
0.7	0.103957936	27.97610	0.3184735	7.917651
0.7	0.240156321	27.95104	0.3189454	7.835833
0.7	0.554792263	27.94103	0.3187177	7.678114
0.7	1.281642117	27.95454	0.3181156	7.369797
0.7	2.960759592	28.09005	0.3163118	6.846891
0.7	6.839738840	28.85612	0.3001402	7.277639
0.7	15.800684229	30.85446	0.2909322	8.878579
0.8	0.008432348	28.01238	0.3175718	8.070449
0.8	0.019479818	28.00936	0.3175163	8.039224
0.8	0.045000908	28.00020	0.3176370	7.972411
0.8	0.103957936	27.97254	0.3185722	7.906252
0.8	0.240156321	27.94734	0.3189795	7.817664
0.8	0.554792263	27.94076	0.3186304	7.640068
0.8	1.281642117	27.96246	0.3178774	7.295586
0.8	2.960759592	28.13781	0.3151418	6.755044
0.8	6.839738840	29.02499	0.2946305	7.449163
0.8	15.800684229	31.33025	0.2909322	9.216728
0.9	0.008432348	28.01155	0.3175817	8.069278
0.9	0.019479818	28.01019	0.3174714	8.033794
0.9	0.045000908	27.99737	0.3177359	7.963507
0.9	0.103957936	27.96951	0.3186458	7.895521
0.9	0.240156321	27.94458	0.3189922	7.800738
0.9	0.554792263	27.93743	0.3186621	7.602655
0.9	1.281642117	27.97328	0.3175364	7.223312
0.9	2.960759592	28.19981	0.3131783	6.691255
0.9	6.839738840	29.16270	0.2914991	7.572163
0.9	15.800684229	31.88869	0.2909322	9.602631
1.0	0.008432348	28.00984	0.3176179	8.064194
1.0	0.019479818	28.01048	0.3174197	8.027480
1.0	0.045000908	27.99330	0.3178932	7.955851
1.0	0.103957936	27.96647	0.3187201	7.885657
1.0	0.240156321	27.94291	0.3189672	7.784590
1.0	0.554792263	27.93467	0.3186856	7.565986
1.0	1.281642117	27.98635	0.3171131	7.153438
1.0	2.960759592	28.26663	0.3108182	6.661504
1.0	6.839738840	29.27235	0.2909322	7.666803
1.0	15.800684229	32.54609	0.2909322	10.044511

RMSE was used to select the optimal model using the smallest value.

The final values used for the model were  $\alpha = 1$  and  $\lambda = 0.5547923$ .

```
>
> #create a graph displaying the accuracy of all models
> plot(model)
> plot(varImp(ridge$finalModel))
> plot(cv)
> plot(ridge)
> hist(Features$X0.19,col = "green")
> hist(Features$X24,col = "red")
> hist(Features$X11.291044776119403,col = 'yellow')
> fit = glmnet(x, y)
> plot(fit)
> cvfit = cv.glmnet(x, y)
> plot(cvfit)
> tfit=glmnet(x,y,lower=-.7,upper=.5)
> plot(tfit)
```

#compare with linear models and report the accuracy

```
cor(Features_Variant_1$X0.19,Features_Variant_1$X24)
```

```
mod=lm(Features_Variant_1$X0.19~Features_Variant_1$X1)
```

```
predict(mod)
```

```
Features_Variant_1$error=mod$residuals
```

```
library(car)
```

```
dwt(mod)
```

```
plot(Features_Train$X0.19,Features_Train$X24,
abline(lm(Features_Variant_1$X0.19~Features_Variant_1$X1),col='red'))
```

```
#Assumption1 Linearity
```

```
plot(Features_Variant_1$X0.19,Features_Variant_1$error, xlab="X24",ylab="Residuals",
main="Linearity")
```

```
#Assumption - Normality
```

```
hist(Features_Variant_1$error, xlab = "Residuals",main= "Histogram of Residuals", col="yellow")
```

```
#Running Regression
```

```
fit<-lm(X0.19~X24+X463+X11.291044776119403+X1.0+X70.49513846124168,  
data=Features_Variant_1)
```

```
fit
```

**#Prediction Accuracy- the one which has good prediction accuracy; in other words, which has the smallest prediction error. Consider the simple case of fitting a linear regression model to the observed data. #A model is a good fit, if it provides a high R2 value.**

**#Coefficients, Significance of slope, R Square, Model Fit**

```
summary(fit)
```

```
#Multicollinearity
```

```
vif(fit)
```

```
> cor(Features_Variant_1$X0.19,Features_Variant_1$X24)  
[1] 0.01258501  
> mod=lm(Features_Variant_1$X0.19~Features_Variant_1$X1)  
> summary(mod)  
  
Call:  
lm(formula = Features_Variant_1$X0.19 ~ Features_Variant_1$X1)  
  
Residuals:  
    Min       1Q   Median       3Q      Max   
-10.37   -8.27   -6.32   -3.03  1295.68   
  
Coefficients:  
              Estimate Std. Error t value Pr(>|t|)      
(Intercept)    10.502642    0.275383   38.14   <2e-16 ***  
Features_Variant_1$X1 -0.131088    0.008768  -14.95   <2e-16 ***  
---  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
  
Residual standard error: 35.4 on 40946 degrees of freedom  
Multiple R-squared:  0.005429, Adjusted R-squared:  0.005404  
F-statistic: 223.5 on 1 and 40946 DF, p-value: < 2.2e-16
```

```
> predict(mod)
```

```
> Features_Variant_1$error=mod$residuals  
> library(car)  
Loading required package: carData
```



```
Attaching package: 'car'
```

```
The following object is masked from 'package:dplyr':
```

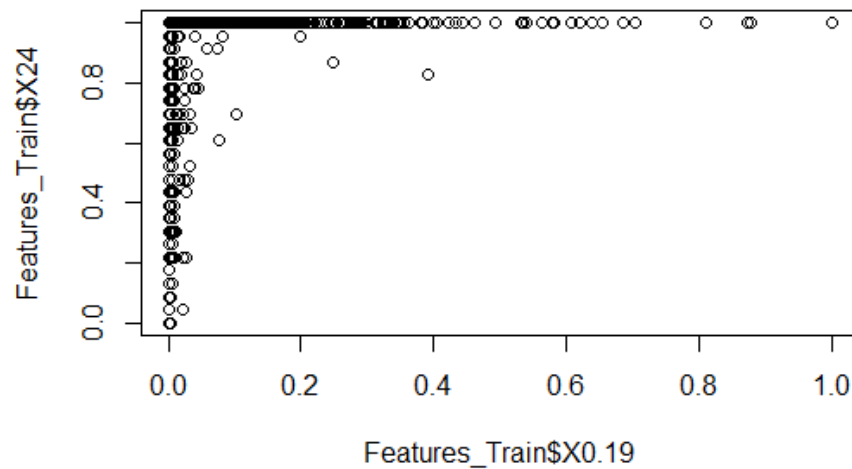
```
  recode
```

```
The following object is masked from 'package:purrr':
```

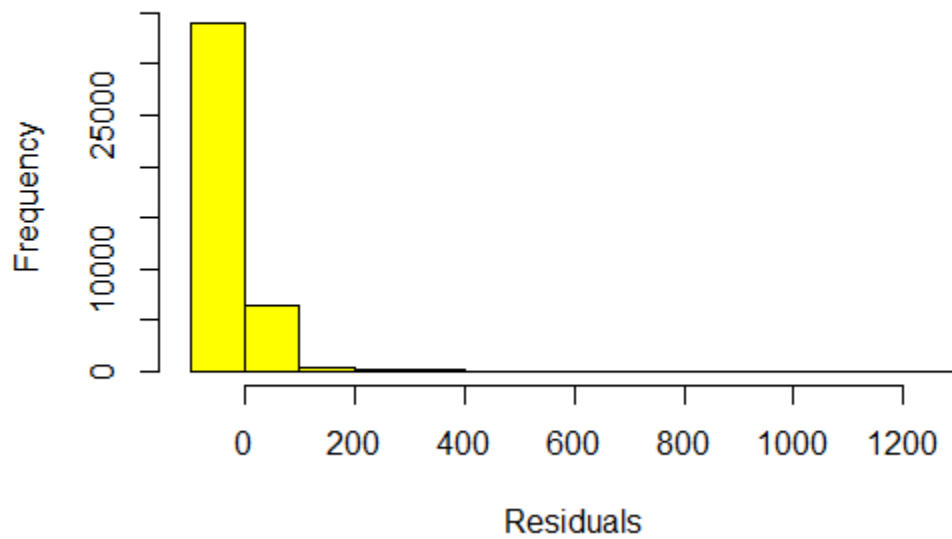
```
  some
```

```
> dwt(mod)
lag Autocorrelation D-W Statistic p-value
1      0.125323      1.749352      0
Alternative hypothesis: rho != 0
```


```
> plot(Features_Train$X0.19, Features_Train$X24, abline(lm(Features_Variant_1$X0.19 ~ Features_Variant_1$X1), col='red'))
> hist(Features_Variant_1$error, xlab = "Residuals", main = "Histogram of Residuals", col="yellow")
> plot(Features_Variant_1$X0.19, Features_Variant_1$error, xlab="X24", ylab="Residuals", main="Linearity")
```



### Histogram of Residuals



### Linearity



```
> fit<-lm(x0.19~x24+x463+x11.291044776119403+x1.0+x70.49513846124168, data=Fe  
atures_Variant_1)  
> fit<-lm(x0.19~x24+x463+x11.291044776119403+x1.0+x70.49513846124168, data=Fe  
atures_Variant_1)  
> fit
```

```
Call:
lm(formula = x0.19 ~ x24 + x463 + x11.291044776119403 + x1.0 +
    x70.49513846124168, data = Features_Variant_1)

Coefficients:
(Intercept)                x24                x463  x11.2910447761
19403
-1.765e+01                7.229e-01                3.037e-06                6.93
0e-02
                x1.0  x70.49513846124168
                5.878e-02                2.513e-02
```

```
> summary(fit)

Call:
lm(formula = x0.19 ~ x24 + x463 + x11.291044776119403 + x1.0 +
    x70.49513846124168, data = Features_Variant_1)

Residuals:
    Min       1Q   Median       3Q      Max
-235.02   -5.38   -1.03    0.17  1266.57

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -1.765e+01  2.079e+00  -8.490  < 2e-16 ***
x24           7.229e-01  8.654e-02   8.354  < 2e-16 ***
x463          3.037e-06  1.785e-06   1.701  0.088916 .
x11.291044776119403  6.930e-02  1.775e-02   3.905  9.46e-05 ***
x1.0          5.878e-02  1.530e-02   3.841  0.000123 ***
x70.49513846124168  2.513e-02  7.821e-03   3.213  0.001315 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 33.41 on 40942 degrees of freedom
Multiple R-squared:  0.1141, Adjusted R-squared:  0.114
F-statistic: 1054 on 5 and 40942 DF, p-value: < 2.2e-16
```

```
> vif(fit)

                x24                x463  x11.291044776119403                x1
.0
1.012425                1.438330                87.332985                42.0345
25
x70.49513846124168
14.928307
```

```
plot(cv$lambda.min)

plot(model)

plot(model$bestTun)

plot(model)
```

```
plot(varImp(ridge$finalModel))
```

```
plot(cv)
```

```
plot(ridge)
```

```
hist(Features$X0.19,col = "green")
```

```
hist(Features$X24,col = "red")
```

```
hist(Features$X11.291044776119403,col = 'yellow')
```

```
fit = glmnet(x, y)
```


```
plot(fit)
```


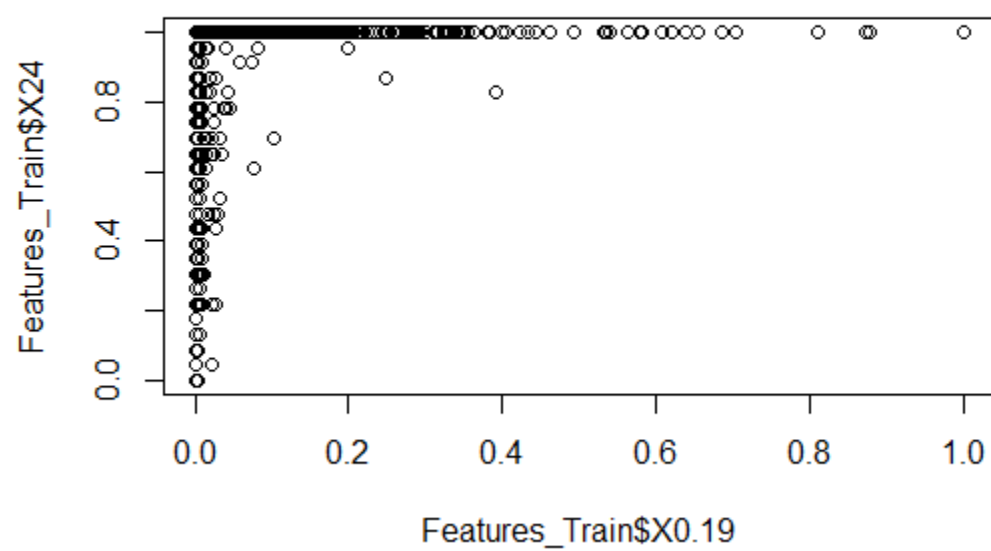
```
cvfit = cv.glmnet(x, y)
```



```
plot(cvfit)
```

```
tfit=glmnet(x,y,lower=-.7,upper=.5)
```


```
plot(tfit)
```








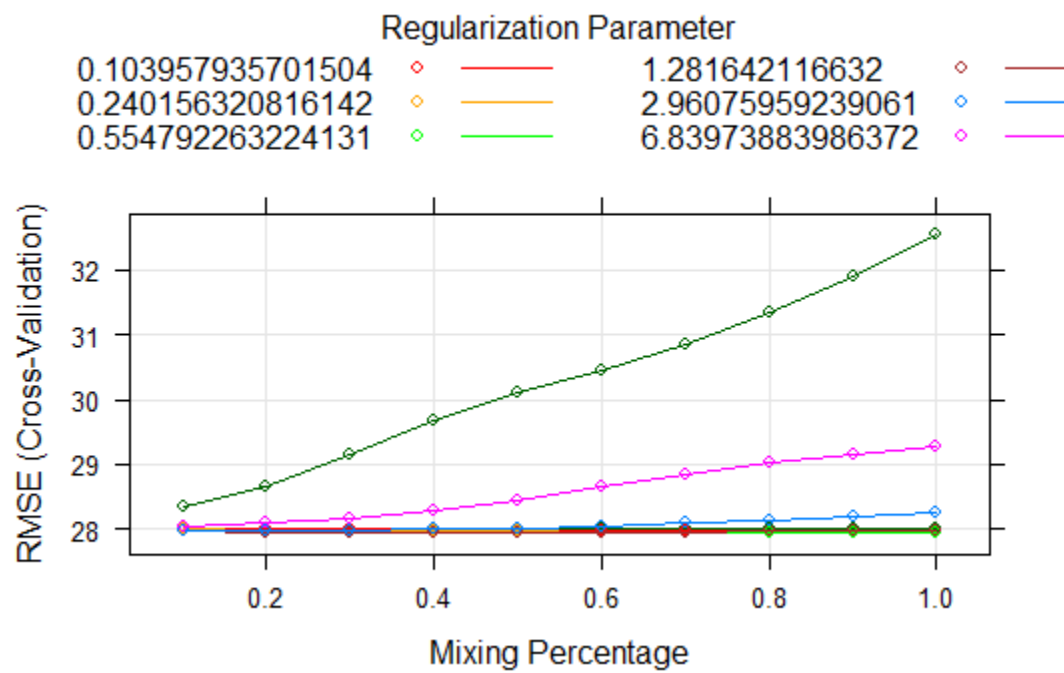
plot( lasso)




plot( predictions)



```
plot(model)
```




```
plot(varImp(ridge$finalModel))
```

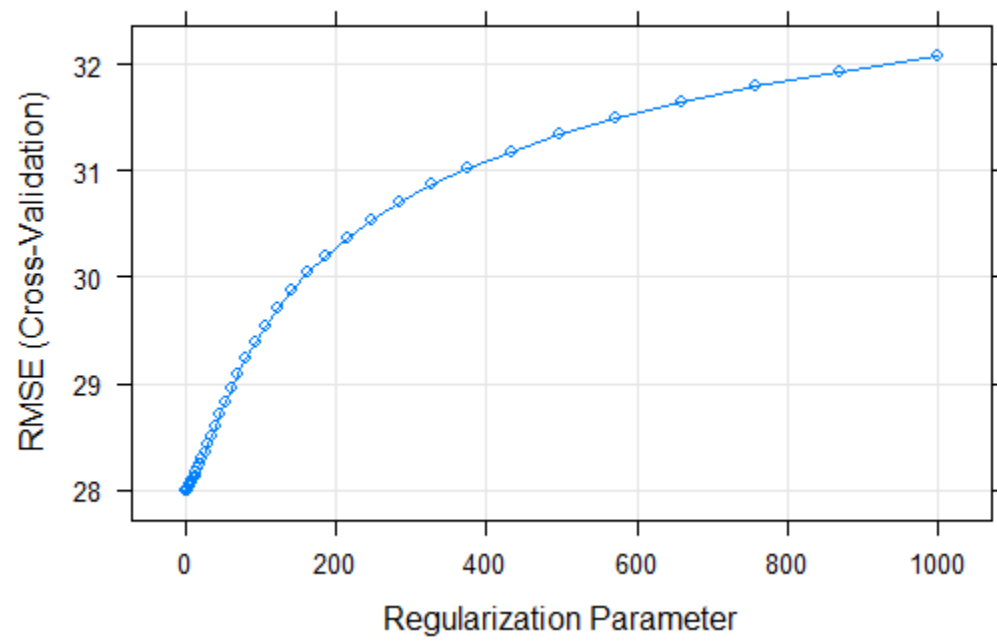





plot(cv)



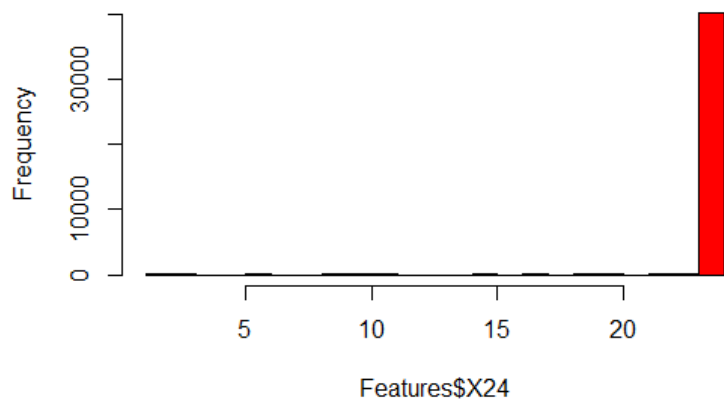
plot(ridge)




**Histogram of Features\$X0.19**




**Histogram of Features\$X24**



**Histogram of Features\$X11.291044776119403**




**Hits@10**



```
fit = glmnet(x, y)
```

```
plot(fit)
```



```
tfit=glmnet(x,y,lower=-.7,upper=.5)
```

```
plot(tfit)
```

